

TRINITY COLLEGE · SEPTEMBER 14, 2026

AI Integration

Your first coding agent

Ken Kousen · CPSC 415

Monday 1:30–4:10 · MC-225

An idea becomes something you can use

Help me memorize a text.

One line at a time.

A Next button.

Start over after the last line.

Watch the files change. Then try the result.

What you bring to the work

The agent

- Reads files
- Writes code
- Runs available tools
- Proposes changes


You

- Define the outcome
- Choose and refine
- Check what happened
- Explain the result



A working page is the beginning of the conversation.

A small start, increasing responsibility



Doghouse

Make an idea tangible.

Cabin

Make it dependable.

Model, harness, application

Part	Today
Model	The model selected through OpenRouter
Harness	Claude Code: context, tools, permissions, loop
Application	Your HTML page running in a browser

The agent uses AI to build the app.

The app itself makes no AI calls today.

The agent works in a loop

Read → propose → act → observe

Then decide what to do next.



You can interrupt, correct the goal,
or ask it to explain.

Text becomes tokens

Welcome to Trinity Collge. Let's learn something about AI

TOKENS	CHARACTERS
12	57

Welcome to Trinity Collge. Let's learn something about AI

Words, word fragments, and punctuation.
Different tokenizers can split the same text differently.

[The Tokenizer Playground](#)

Know where the bill comes from

OpenRouter

Your account
Your selected model
Your key's spending cap
Your usage record

Today's first candidate

`minimax/minimax-m3`

Pay-as-you-go API usage.
Use the model announced in class.

If access fails, ask. Do not buy a subscription to fix it.

Cost per accepted result

Calls + retries + corrections
+ your time



Measure the work that reaches an acceptable result.

A free endpoint still costs time when it stalls.

A higher price still needs to earn its place.

2:00–2:45 · GUIDED SETUP

Open the Week 1 guide

github.com/kousen/ai-integration-course

[weeks / 01](#)

Choose **macOS Terminal** or **Windows PowerShell**.
Run one step at a time. Show the error if a step fails.

Three setup checkpoints

Checkpoint	What you should see
A · Installed	Git and Claude Code print a version
B · Project ready	Your own folder, poem, and Git branch
C · Connected	Agent reads 14 lines; correct model in usage

Tell me the checkpoint where you are stuck.

Keep two folders separate

```
cpsc415/  
  ai-integration-course/  ← guides and launcher  
  lyrics-trainer/         ← your files and commits  
    lyrics.txt  
    .gitignore
```



Start the agent inside `lyrics-trainer`.

Launch through your wrapper

macOS Terminal

```
bash ../ai-integration-course/scripts/orclaude "minimax/minimax-m3"
```



Windows PowerShell

```
& ..\ai-integration-course\scripts\orclaude.bat "minimax/minimax-m3"
```



First enter your key using the hidden prompt in the guide.
The instructor may substitute the rehearsed fallback model.

Verify the connection

Read `lyrics.txt`.

How many nonblank lines does it contain?

Do not change any files yet.

1. Check the answer against the file: **14 lines**.
2. Inspect `/status`.
3. Check the model in OpenRouter Activity.

The agent's self-description is not billing evidence.

Use only what this task needs

- Work inside your exercise folder.
- Use the supplied poem, not someone else's personal data.
- Keep the API key out of files, screenshots, and prompts.
- Read a proposed action before approving it.



A model can suggest something you should decline.

Build the first version

Read the supplied poem.

Show one line and a counter.

Next advances; the last line wraps to the first.

One HTML file. No packages. No API calls.

Show a short plan and wait for approval.

The complete prompt is in [the lab handout](#).

Choose one improvement

Previous · Restart · Keyboard shortcuts
Readability



Decide the behavior before asking for the change.

Previous on line 1: stay there, or wrap to line 14?

A useful change request

Add a Restart button.

It returns to line 1 and updates the counter.

Keep Next and wraparound working.

Keep the app in one file.

Explain the change before editing.

What stays the same is part of the requirement.

Check the boundaries

Action	Expected result
Open the page	First line; counter 1 of 14
Next once	Second line; counter 2 of 14
Next from the last line	First line; counter 1 of 14
Keyboard activation	The focused button works
Your change at an edge	The behavior you specified

Record the actual result in `CHECKS.md`.

When a check fails

Report

What you did
What you expected
What you observed

Verify the correction

Read the explanation
Try the failing case again
Check related behavior



If nothing fails, add a check. Do not invent a failure.

Explain one piece of the code

Where is the current line stored?

What happens after the last line?

How does the counter change?

| Point to the implementation, then show its behavior.

Save evidence of the work

<code>index.html</code>	The app with your improvement
<code>lyrics.txt</code>	The supplied text
<code>CHECKS.md</code>	Expected and observed behavior
<code>README.md</code>	How to run it; choices and learning
<code>.gitignore</code>	Keep local credentials out



Commit the work. Tag it `week01-submitted`.

The rest of the semester

Component	Weight
Participation and labs	15%
Academy certificates	10%
Team projects: multimodal / agentic-MCP	15% / 20%
Individual AI-enabled profile site	25%
Individual Comprehension Defense	15%

Choose and justify your project languages.
One guided exercise will explore an unfamiliar language.

The process grows with the work

Week	New evidence
1	Working app, checks, explanation, commits
2–3	Intent, then specification
4–5	Written plan, then automated checks
6 onward	Reviewed pull requests and the full loop

Write down decisions while they are being made.

Before next Monday

September 21 · 1:30 PM

- Submit your repository URL and `week01-submitted` tag.
- Upload your **Claude Code 101** certificate.
- Read the syllabus and submit the **Academic Honesty Pledge**.

Start now: AI Fluency for Students, due September 28.

Week 1 guides and public examples are linked from Moodle.
Selected draft book readings will be shared there separately.

Before you leave

What did you ask for?

What did you change?

What did you check?



Show the result, and one thing you understand about it.

Three interfaces, one application

Compare. Choose. Refine.

Then check the behavior again.

Prof. Kousen demonstrates Anthropic's `/design` using his own access.
This command is not part of the OpenRouter lab.

Sources and next steps

- Ken Kousen, *Claude Code: Up and Running*: Chapters 1–4 and selected Chapter 11 material. Draft readings via Moodle.
- [Companion examples and launcher](#)
- [Claude Code training](#) · [Codex training](#)
- [Practical AI Literacy](#): tokenizer screenshots and cost/verification teaching material
- [Tokenizer Playground](#)
- [Claude Code setup](#) · [OpenRouter integration](#)
- Shakespeare, Sonnet 18: public-domain lab text. Project-scale illustration from Prof. Kousen's book, used for this course.

Prepared September 9, 2026. Recheck models and service interfaces before class.